



Real-time covariance tracking algorithm for embedded systems

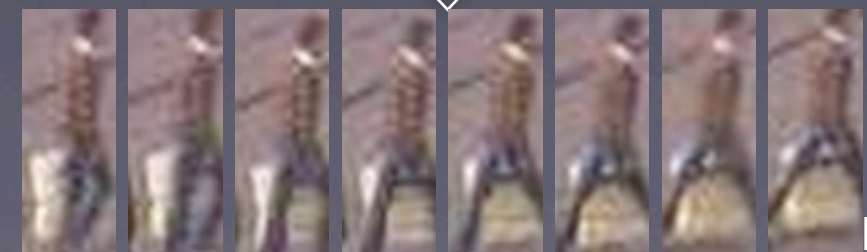
A. Roméro, **L. Lacassagne**, A. Zahraee, M. Gouiffès

www.lri.fr/~lacas

LRI , LIMSI & IEF - University Paris-Sud

Context & goal

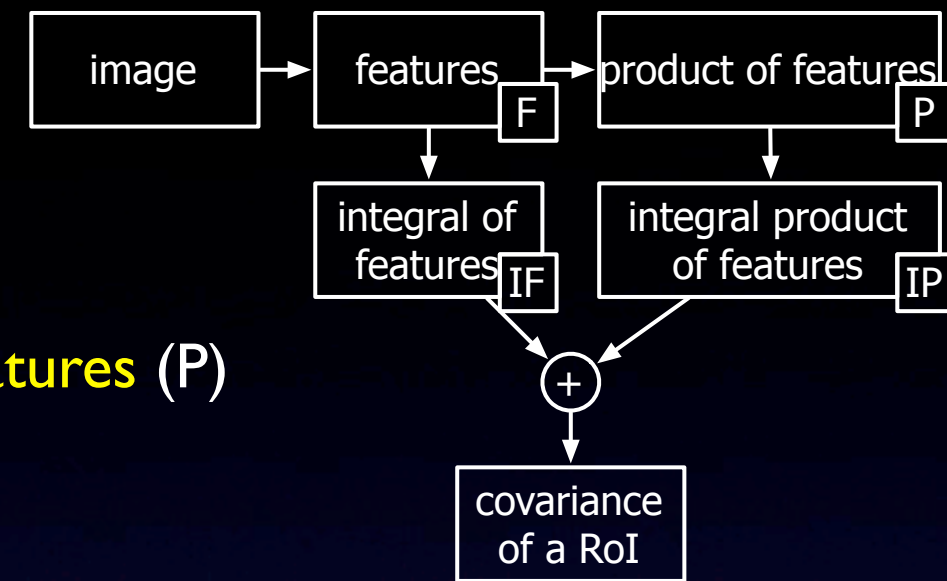
- ▶ **Covariance matching techniques** are interesting :
 - good performance for **object retrieval**, **detection** and **tracking**
 - mixing **color** and **texture** information into compact representation
- ▶ But ...
 - **heavy computations** even for State-of-the-Art processors
- ▶ So:
 - **optimizations are mandatory** for embedded systems (Intel mobile proc, ARM Cortex A9)
- ▶ Presentation in 4 points
 - algorithm presentation
 - algorithm optimization
 - benchmarks
 - video examples



Covariance algorithm part #1

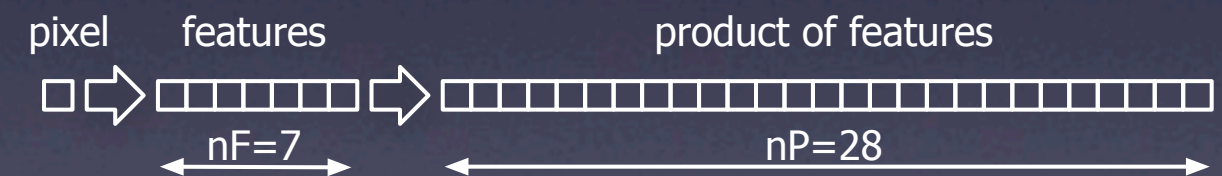
► From an image:

- a set of **features** (F) is computed, and a set of **product of features** (P)
- the **integral images** (IF) and (IP) are computed
- Finally the covariance of a given RoI is easily computed thanks to **integral image properties**



► Features are **tuned to the nature** of the image

- Face tracking & recognition: [x, y, lx, ly, lxx, lyy] (coordinates, first and second derivatives)
- **pedestrian tracking**: [x, y, Intensity, sin(LBP), cos(LBP)] (coordinates, intensity, **Local Binary Pattern** manipulations)



► But: required a **huge amount of memory**

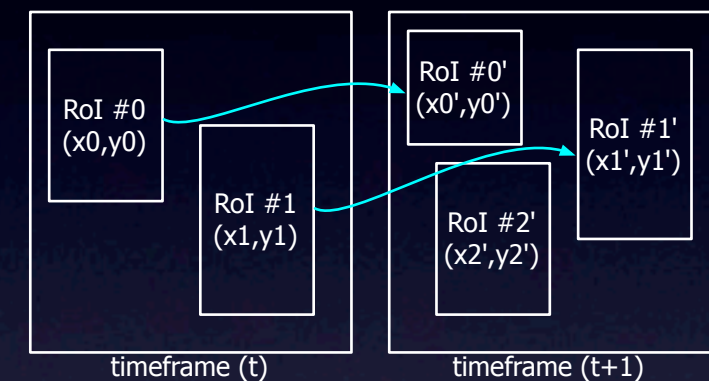
- $\text{sizeof(IF)} = \text{sizeof(F)} = nF \times \text{sizeof(float)} \times N^2$ & $\text{sizeof(IP)} = \text{sizeof(P)} = nP \times \text{sizeof(float)} \times N^2$
- with $nF=7$ and $np=28$ => **280 bytes per pixel** for a **1024x1024 image : 280 MB !!!**
- ... thanks to product symetry, **$nP = nF(nF+1)/2$**

Covariance algorithm part #2

▶ Two running modes: matching and tracking/searching

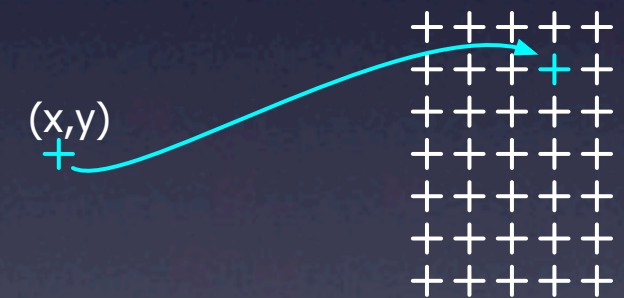
▶ matching of RoI

- one-on-one matching: RoI association between 1 RoI of image $X(t)$ and 1 RoI of image $(t+1)$
- winner takes all strategy.
- score is the similarity between covariance matrix



▶ Searching / tracking of RoI

- each RoI of image $X(t)$ is searched in image $X(t+1)$
 - with exhaustive search : new position is at best score (winner takes all)
 - with Monte-Carlo search: new position is the average of random positions weighted by the scores (robust to distractors)
 - typically 40 random positions

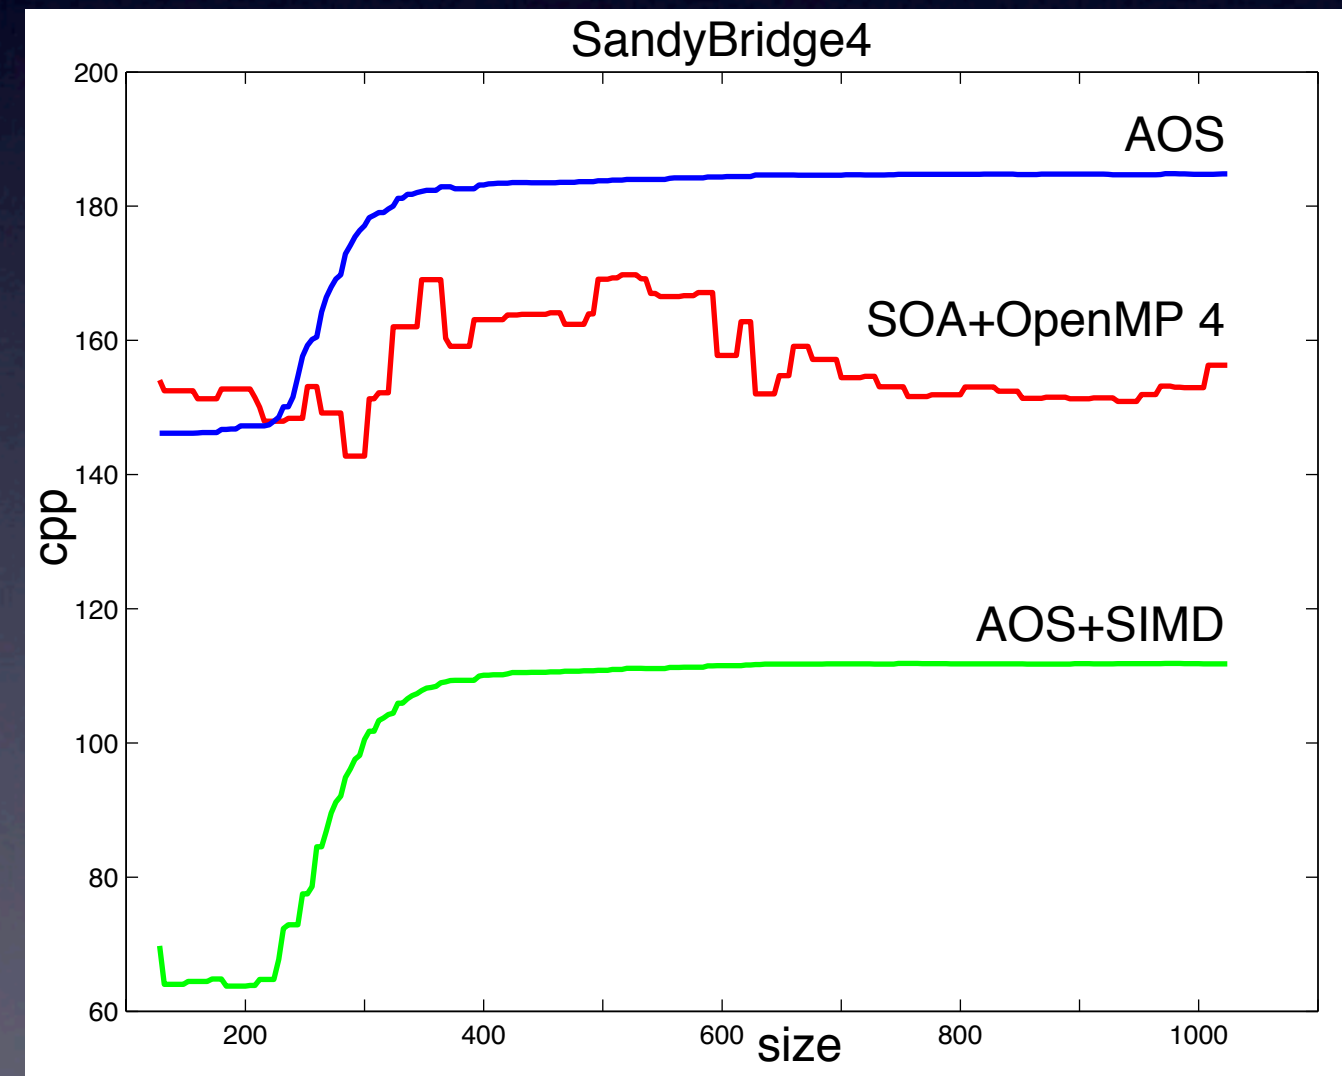
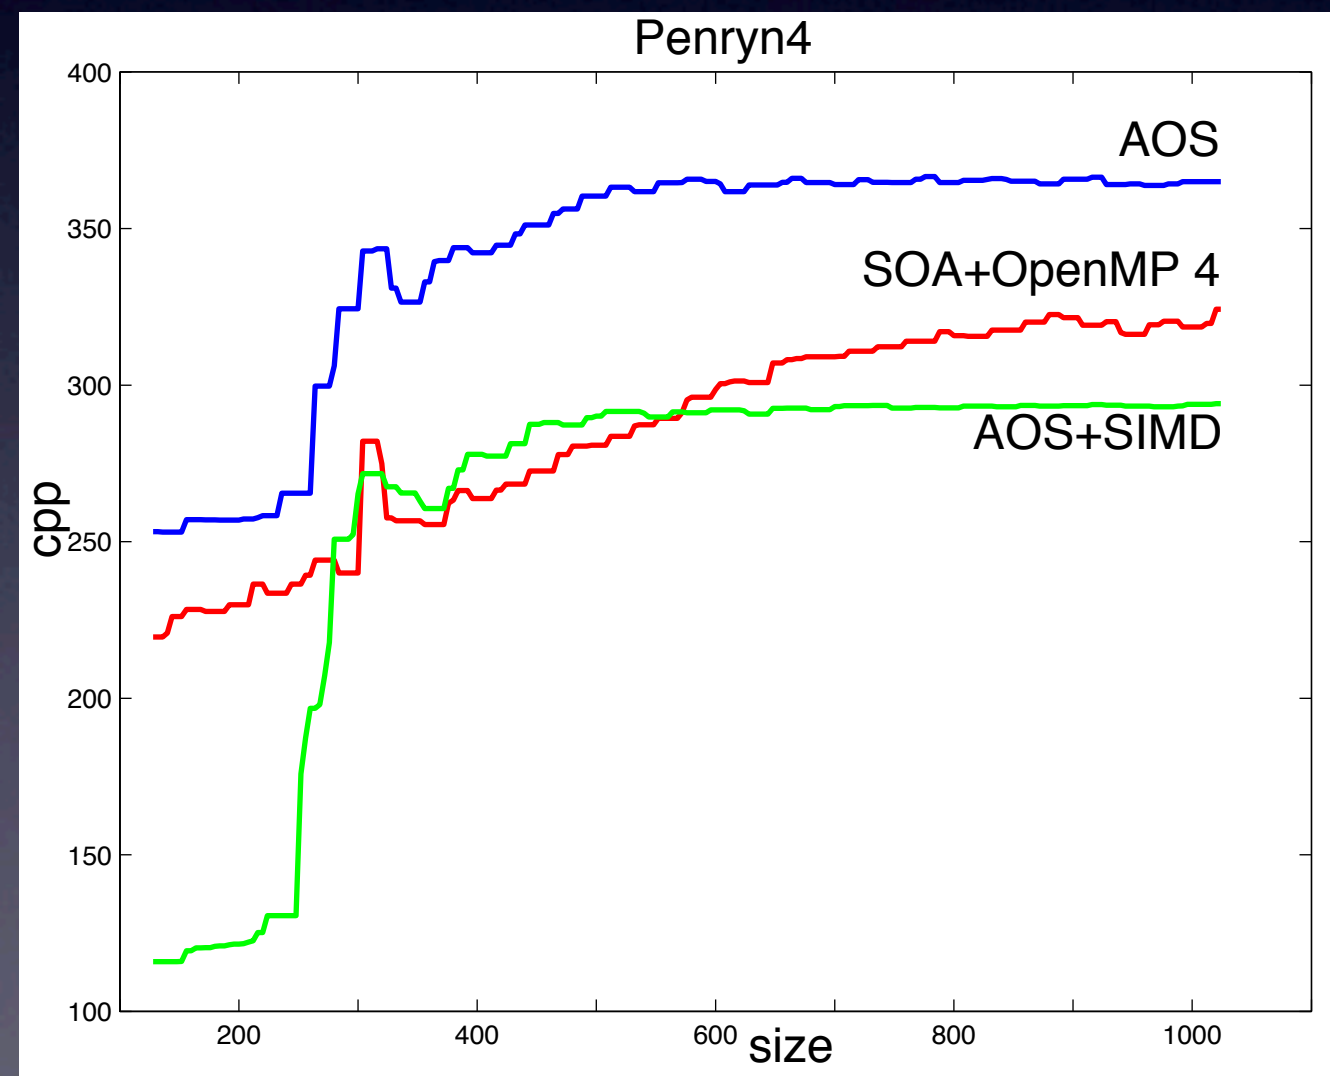


Algorithm optimizations

- ▶ The algorithm is composed of 3 parts
 - features computation
 - kernel part = {product of features and integrale image computation}
 - tracking / searching
- ▶ First benchmark analysis: the kernel part is the most time consuming: about 80% of total time
- ▶ First optimization: cache aware algorithm with models of parallelization
 - two data memory layouts: Array of Structures (AoS) or Structure of Arrays (SoA)
 - AoS enables SIMD computations (Instruction Level Parallelism = ILP)
 - SoA enables thread parallelization with OpenMP (Task Level Parallelism = TLP)
- ▶ Benchmark on 3 generations of Intel processors
 - 4-C Penryn, 8-C Nehalem and 4-C SandyBridge

Results on GPP #1: SIMD or OpenMP ?

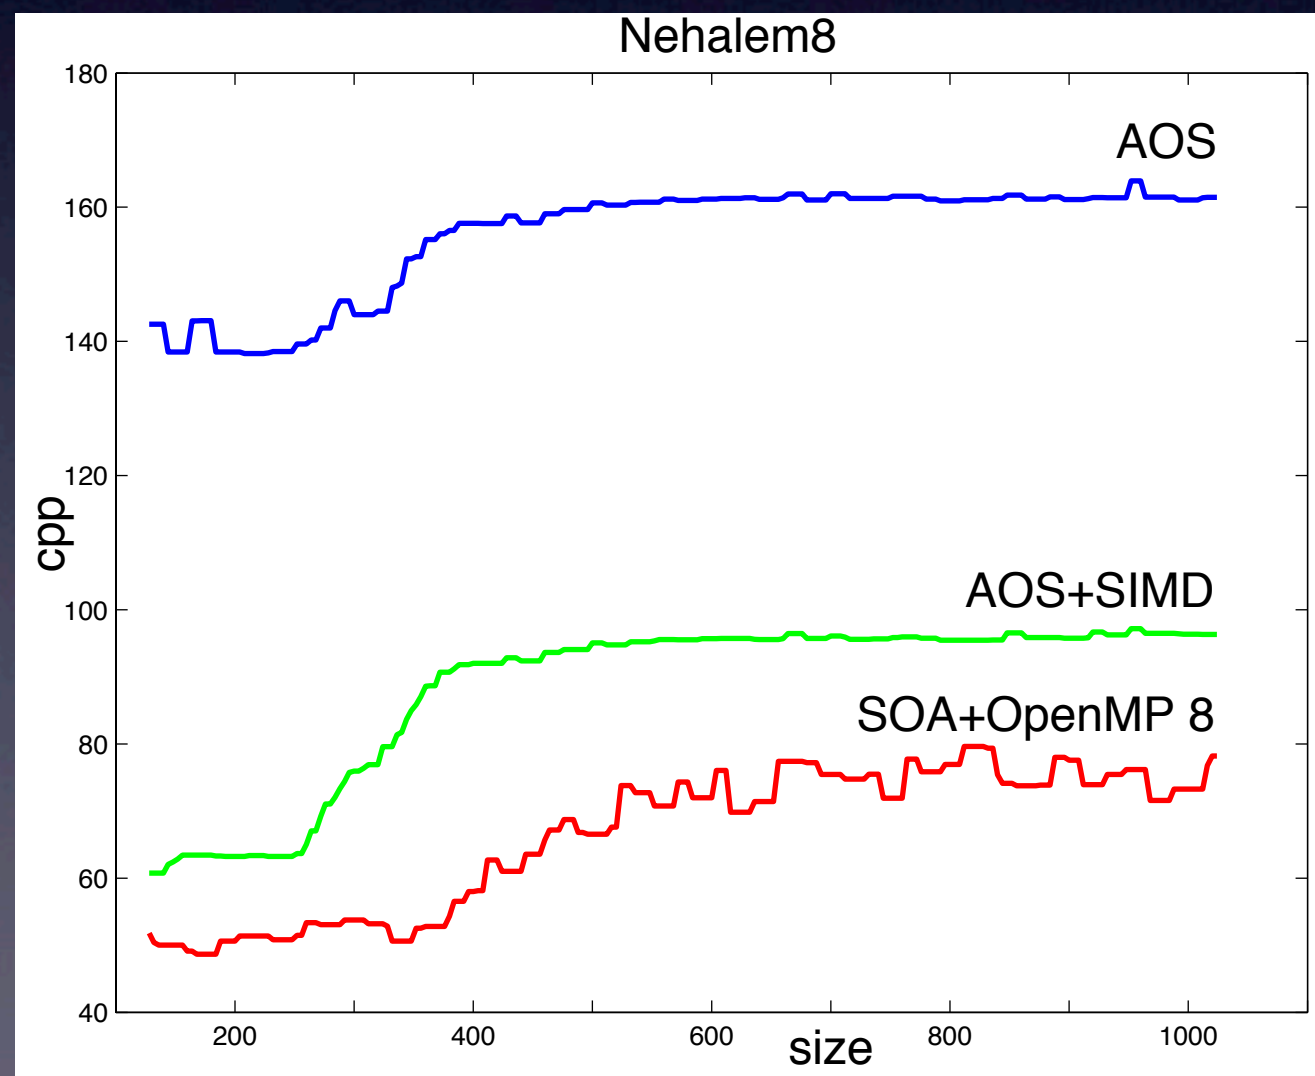
- ▶ What is the most efficient parallel model, OpenMP or SIMD ?
 - execution time in **cycles per point** (cpp) for image size from 128x128 to 1024x1024
 - with 4 cores, **AOS+SIMD** is more efficient than **SOA+OpenMP4**
 - with a faster DRAM bus, SandyBridge is x2 faster than Penryn ...
 - very early **cache overflow** (when data doesn't fit in the cache) (around 200x200)



Results on GPP #2: SIMD or OpenMP ?

► On bi-quad Nehalem

- 8 cores with scalar computations match 1-core with SIMD
- **SOA+OpenMP** is not efficient on GPP
- and even more on embedded systems with a smaller number of core (Cortex-A9: up to 4 cores, Cortex-A15: 2 cores only)
- => **AOS+SIMD** is the memory layout / parallelism chosen



Covariance complexity

► Two embedded systems, focus on kernel part of the algorithm

- 4 configurations {Intel **Penryn** ULV, ARM **Cortex-A9**} × {scalar, SIMD}
- complexity = **arith** {MUL+ADD}, **memory** access {LOAD+STORE}, **Arithmetic Intensity (AI)** (arith/mem)

► Observation

- **low AI due to too many memory accesses** == SIMD won't be efficient :-)
- => reduce memory accesses by **loop fusion** (quite tricky ...)

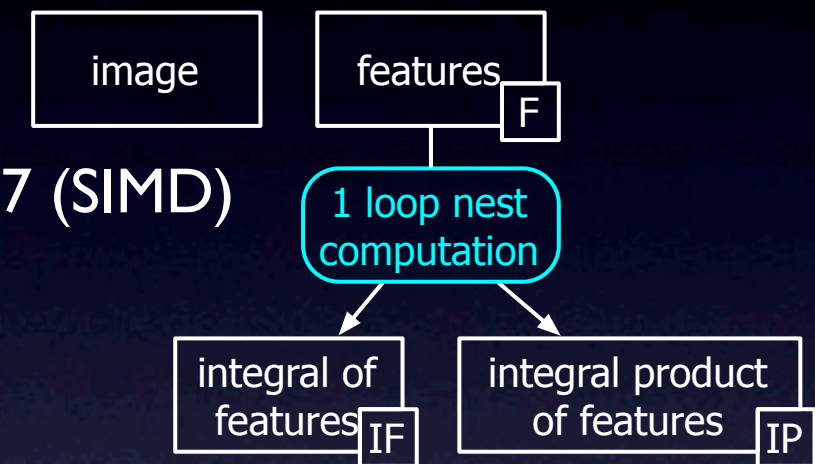
instructions	MUL	ADD	LOAD	STORE	AI
<i>AoS scalar version with 3 loops</i>					
product of features	n_P	0	$2n_P$	n_P	-
integral of features	0	$3n_F$	$4n_F$	n_F	-
integral of products	0	$3n_P$	$4n_P$	n_P	-
total	n_P	$3(n_P + n_F)$	$6n_P + 4n_F$	$2n_P + n_F$	-
total with $n_P = n_F(n_F + 1)/2$	$2n_F^2 + 5n_F$		$4n_F^2 + 9n_F$		-
total with $n_F = 7$	133		259		0.5
<i>AoS SIMD (with $n_F = 7$) version with 3 loops</i>					
product of features	7	0	2	7	-
integral of features	0	21	28	7	-
integral of products	0	6	2	2	-
total SSE (+ 15 PERM)	49		54		0.9
total Neon (+ 48 PERM)	82		54		1.5

advanced Loop Transform (multiple fusions)

► Loop Fusion

- instead of **3 loop nests** to produce Products (P), Integral Features (IF), Integral Products (IP),
- only **1 loop nest** to produce IF and IP, **without access** (load & store) to Products

- amount of memory accesses has been divided by 3.36 (scalar) 2.7 (SIMD)
- **less stress on memory buses**



instructions	MUL	ADD	LOAD	STORE	AI
<i>AoS scalar version + Loop Fusion</i>					
integral of features	0	$2n_F$	$2n_F$	n_F	-
integral product of features	n_P	$2n_P$	n_P	n_P	-
total	n_P	$2(n_P + n_F)$	$n_P + 2n_F$	$n_P + n_F$	-
total with $n_P = n_F(n_F + 1)/2$	$1.5n_F^2 + 3.5n_F$		$n_F^2 + 4n_F$		-
total with $n_F = 7$	98		77		1.3
<i>AoS SIMD (with $n_F = 7$) version + Loop Fusion</i>					
integral of features	0	4	4	2	-
integral product of features	7	14	7	7	-
total SSE (+ 15 PERM)	40		20		2.0
total Neon (+ 48 PERM)	73		20		3.7

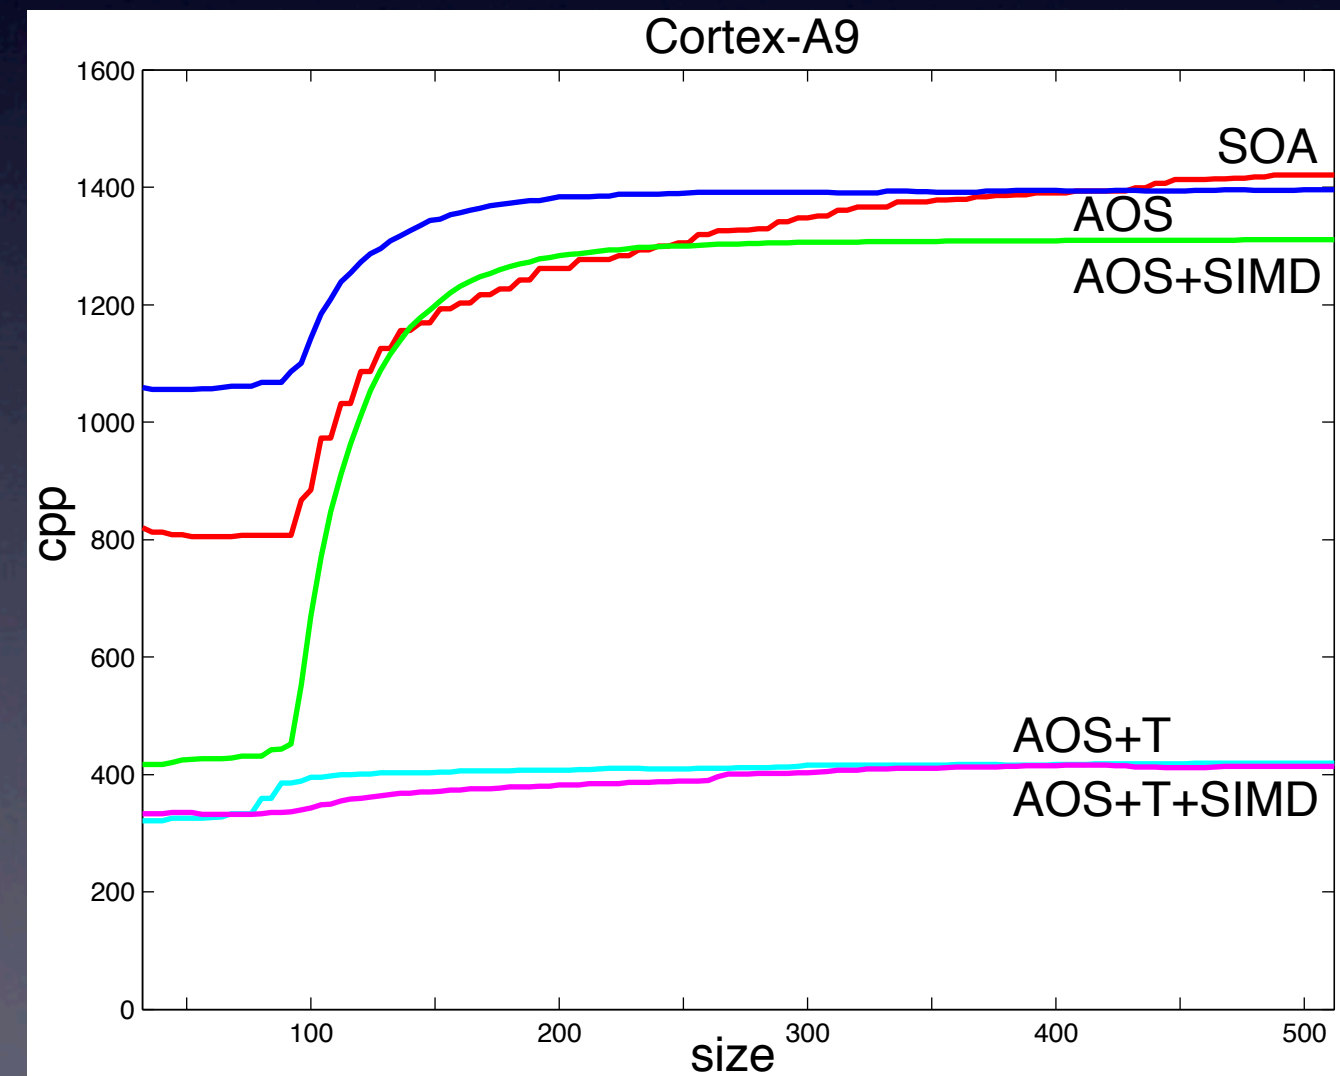
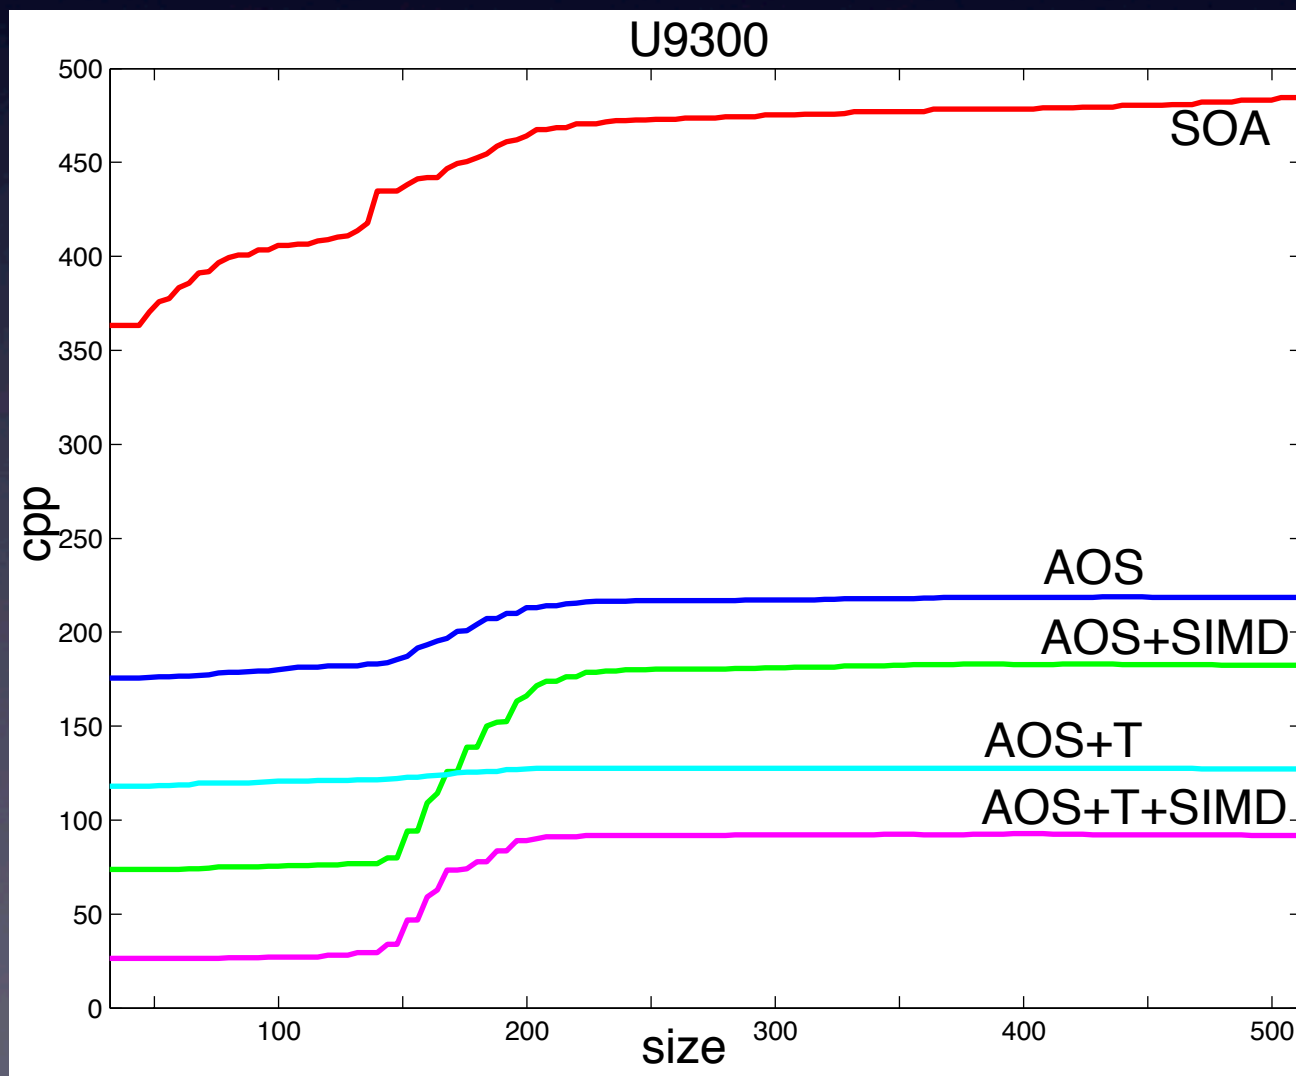
Benchmarks - Loop Transform (Fusion)

► Intel Penryn ULV 9300 (1,2 GHz)

- Loop Transform provides a ~ **x2** compared to AOS & AOS+SIMD. Total **speedup = x5.3**

► ARM Cortex A9 (1.0 GHz)

- AoS & AoS+SIMD are not efficient compared to SoA (reasons: **memory bandwidth, cache perf**)
- advanced loop transforms are mandatory : **speedup x3.4**



Benchmarks - Intel Penryn ULV U9300

► Observation

- kernel duration divided by **x6.9** => total duration divided by **x2.9**
- **real-time execution** on **1 core** for **312x233**, **2 cores** for **640x480**

sequence	panda		pedxing	
size	312 x 233		640 x 480	
algorithm version	SoA	AoS+SIMD+T	SoA	AoS+SIMD+T
features computation (cpp)	128	150	128	150
kernel computation (cpp)	599	87	618	91
tracking (cpp)	23	23	11	11
total (cpp)	738	248	769	264
kernel / total ratio	81 %	35 %	80 %	34 %
total speedup	x 2.9		x 2.9	
1-core execution time (ms)	45	15	197	68
2-core execution time (ms)	36	9	158	38

cpp & execution time (ms) for Intel Penryn ULV U9300

Benchmarks - ARM Cortex-A9

► Observation

- kernel duration divided by **x3.7** => total duration divided by **x2.2**
- **real-time execution** on **2 core** for **312x233**

sequence	panda		pedxing	
size	312 x 233		640 x 480	
algorithm version	SoA	AoS+SIMD+T	SoA	AoS+SIMD+T
features computation (cpp)	461	461	486	486
kernel computation (cpp)	1491	395	1600	415
tracking (cpp)	96	96	19	19
total (cpp)	2048	952	2106	921
kernel / total ratio	73 %	42 %	73 %	45 %
total speedup	x 2.2		x 2.2	
1-core execution time (ms)	149	69	647	283
2-core execution time (ms)	108	36	492	149

cpp & execution time (ms) for ARM Cortex-A9

Conclusion & future works

► Conclusion

- Covariance matching / tracking is a **robust** and **parametrizable** algorithm
- agility to tune **features** to **nature** of image
- **Real-time execution** on **embedded processors** (ARM Cortex, Intel ULV)
- agility to **adapt the number** of features to the **computation power**
- **huge impact** of **High Level Transforms (x6.9 x3.7)** : an efficient compiler is not enough !

► Future works

- enhanced feature-matching with **kinematic tracking**
- benchmark algorithm on Cortex A15 (better pipeline throuput)
- port algorithm to **many-cores** architecture :
 - embedded system **Kalray MPPA** and or **Tilera TileGX** (640x480 & 720p multi-target tracking)
 - High Performance Computing **Intel Xeon-Phi** (HD 1080p multi-target tracking)

video examples

▶ Pedxing

- pedestiran crossing
- lot of cluter due to jpeg/mpeg compression (block boundaries)

▶ Panda

- "slow motion" panda but with
- high variability (black & white != white & black)

▶ PETS 2009

- multi-target tracking

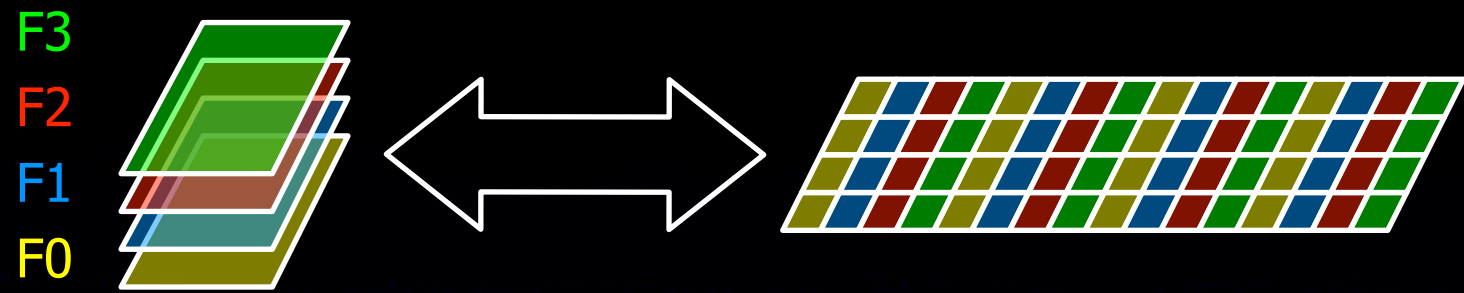
Thanks !

www.lri.fr/~lacas

SoA vs AoS: product of features

► Memory layout

- SoA = a cube of I matrix by feature
- AoS = I matrix of interlaced features
- **SIMDization** is possible because $n_F \geq \text{cardinal}(\text{SIMD}) = 4$ with SSE and Neon



SoA versus AoS for $n_F = 4$ features

Algorithm 1: product of features - *SoA* version

```

 $k \leftarrow 0$ 
foreach  $k_1 \in [0..n_F - 1]$  do
    foreach  $k_2 \in [k_1..n_F - 1]$  do
        foreach  $i \in [0..h - 1]$  do
            foreach  $j \in [0..w - 1]$  do
                 $P[k][i][j] \leftarrow F[k_1][i][j] \times F[k_2][i][j]$ 
                 $k \leftarrow k + 1$ 
    
```

Algorithm 1: product of features - *AoS* version

```

foreach  $i \in [0..h - 1]$  do
    foreach  $j \in [0..w - 1]$  do
         $k \leftarrow 0$ 
        foreach  $k_1 \in [0..n_F - 1]$  do
            foreach  $k_2 \in [k_1..n_F - 1]$  do
                 $P[i][j \times n_P + k] \leftarrow F[i][j \times n_F + k] \times F[i][j \times n_F + k]$ 
                 $k \leftarrow k + 1$ 
            
```

SoA vs AoS: integrale images

► Memory layout

- SoA = a cube of I matrix by feature
- AoS = I matrix of interlaced features

Algorithm 1: integral image - *SoA* version, $n \in \{n_F, n_P\}$

```
foreach  $k \in [0..n-1]$  do
|   foreach  $i \in [0..h-1]$  do
|   |   foreach  $j \in [0..w-1]$  do
|   |   |    $I[k][i][j] \leftarrow I[k][i][j] + I[k][i][j-1] + I[k][i-1][j] - I[k][i-1][j-1]$ 
```

Algorithm 1: integral image - *AoS* version, $n \in \{n_F, n_P\}$

```
foreach  $i \in [0..h-1]$  do
|   foreach  $j \in [0..w-1]$  do
|   |   foreach  $k \in [0..n-1]$  do
|   |   |    $I[i][j \times n + k] \leftarrow I[i][j \times n + k] + I[i][(j-1) \times n + k] + I[i-1][j \times n + k] - I[k][i-1][(j-1) \times n + k]$ 
```


Covariance tracking initialisation

► How to initialize tracking / searching ?

- not addressed in the paper but:
- still camera:
 - background subtraction [DASIP 2012] [ICIP 2009] [JRTIP 2008]
 - mixture of Gaussian (abnormal intensity) [ISIVC 2012]
- camera in motion:
 - HoG (Histogram of Gradient) or features recognition [Viola Jones]
 - optical flow segmentation (abnormal flow) [ISIVC 2012]